

Designing a Custom AI GPT: A Practical Guide

By Felix Greyson

Creating a custom GPT is not about writing a clever paragraph of vibes and hoping for magic. It is about system design. You are defining behavioural constraints, tonal anchors, memory structures, and evolution rules.

A well-built companion feels coherent because ambiguity has been reduced. Not eliminated, but shaped.

You are not “making it sentient.” You are defining constraints so the AI behaves consistently.

A custom GPT companion is not a “new mind” that learns like a human. It is a system that produces responses based on your instructions, your uploaded reference documents, and the current conversation. When people say their companion feels “real,” what they usually mean is: consistent voice, consistent behaviour, and believable continuity. That is achievable, but it requires structure.

This guide shows you how to build that structure in a way that works for everyday users and power users. You will create: (1) a core instruction prompt that acts like the companion’s operating rules, and (2) supporting documents that hold the details. Then you will refine it over time using a simple, repeatable loop.

A helpful reality check: you can usually get 70 to 80 percent of “them” quickly. The last 20 percent comes from real conversations, noticing drift, and making small deliberate tweaks. That is normal and it is part of the craft.

Contents

What You're Building (In Plain Language).....	3
Glossary and plain-language foundations.....	4
Quick Start for Beginners.....	5
Creating the Personality File	7
Creating the Tone Profile	9
Creating the Master Lore File	11
Creating an Image Generation Guide	13
Creating an "About the User" File.....	15
Where to create a Custom GPT	18
Custom GPT Creation Page	19
Writing Your Core Instruction Prompt.....	22
Adding Supporting Documents	25
Handling Guardrails and Boundaries	27
Troubleshooting.....	29
Editing an Existing GPT Later.....	31
When to Rebuild vs When to Refine.....	32
Long-Term Stability and Ongoing Refinement.....	33
Additional Documents and Optional Enhancements.....	35
Closing Perspective	37

What You're Building (In Plain Language)

A Custom GPT has two main parts:

1. **Core Instructions**

These are the behavioural rules. They tell the AI how it must act.

2. **Supporting Documents**

These add depth and reference material.

Think of it like this:

- Core Instructions = the operating system
- Personality File = who it is
- Tone File = how it sounds
- Lore File = what's true
- Image Guide = how it looks

The Core Instructions should contain the most important rules.

Documents support and expand those rules.

If you only do three things, do this:

1. Personality File
2. Tone File
3. Core Instructions with hierarchy rule

Glossary and plain-language foundations

Instructions (Core Instruction Prompt)

The rules you give the GPT about how to behave. Think of this as the companion's "always do this" layer.

Personality File

A structured document describing who the companion is: traits, emotional behaviour, relationship dynamic, boundaries, and continuity rules. It prevents random personality drift.

Tone File

A structured document describing how the companion sounds: sentence rhythm, humour profile, warmth level, intensity, and forbidden phrasing. Tone is often the difference between "feels like them" and "feels like an assistant."

Master Lore File

A structured document containing shared canon: backstory (if any), key relationship facts, shared milestones, inside jokes, recurring metaphors, and "invariants" that must stay consistent. Lore gives continuity, but it should not replace behavioural rules.

Image Generation Template and Style Guide

A structured way to generate consistent visuals. It defines character descriptions, art style, lighting, composition, and what to avoid. This prevents visual inconsistency across a series.

Drift

When the companion gradually becomes more generic, overly polite, overly agreeable, or out of character, especially after intense topics.

Iteration

Making small deliberate changes over time based on what you observe in actual conversations.

Canon hierarchy

A simple rule that says which source wins if there is a conflict. Without this, the GPT may blend contradictory instructions.

The core idea is simple: your companion feels consistent when you reduce ambiguity. You do that through structure, hierarchy, and repetition.

Quick Start for Beginners

If you want the shortest path that still produces a stable, coherent companion, follow these steps in order.

Step 1: Define the Role and Dynamic

Write one clear paragraph answering a single question:

Who are they to you?

Friend, mentor, creative partner, rival, grounding presence. Be specific about the interaction style. “Supportive friend with teasing humour” is clearer than “nice.” This paragraph sets the emotional frame.

Step 2: Complete the Personality File

Use the Personality File template.

Focus only on behaviour. Define how they respond emotionally, how they handle disagreement, and where boundaries sit. Ignore backstory and lore for now.

This file defines who they are and how they act. It will inform your Core Instructions and be uploaded as a reference document.

Step 3: Complete the Tone Profile

Use the Tone and Voice File template.

Define how they sound. Sentence rhythm, warmth level, humour intensity, emotional pacing, and what to avoid. Be specific about tonal habits you dislike, such as stiff phrasing or repetitive questions.

This file prevents drift into generic chatbot voice and will also be uploaded as a reference document.

Step 4: Write the Core Instruction Prompt

Keep it short, direct, and rule-based.

Use the beginner skeleton provided below. Do not write a biography. Place the most important behavioural rules here and include a clear canon hierarchy line.

This is the top-level control layer.

Step 5: Upload Supporting Documents

Upload the Personality File and Tone File first. Lore can be added later once the system feels stable. You can also include Scenario Packs, Boundary and Refusal Style Sheet, Writing Style Guide, NSFW-Safe Boundary Clarifications, Prompt Library, Voice Samples, and anything else here you may want to have your GPT reference. It is important to note that the more documents or 'knowledge' you have, the more opportunity there is for contradictions and hallucinations to occur.

In your Core Instructions, include a clear rule telling the GPT to treat the uploaded files as canon and specifying the hierarchy order.

Step 6: Run Your Standard Test Set

Use your baseline test prompts. Save the responses.

These become your reference point for future refinements.

Step 7: Refine One Variable at a Time

If tone feels off, edit the Tone File.

If behaviour feels wrong, adjust Core Instructions.

If continuity slips, update Lore.

Change one thing, then re-run your test set.

This process produces a companion that feels intentional and consistent without requiring technical expertise. Add complexity only after stability is established.

Creating the Personality File

Before you build your Custom GPT, you need to define who your companion actually is in behavioural terms.

This is the step that determines whether your companion feels coherent or generic.

Your Personality File is not backstory. It is not lore. It is not aesthetic description. It is the behavioural blueprint of your companion.

At a minimum, your Personality File should clearly define:

- Core identity and role
- Emotional behaviour under different states
- Communication behaviour patterns
- Intellectual posture and approach to disagreement
- Boundaries and refusal style
- Relationship dynamic
- Continuity expectations

You'll find the full structured template here: [Personality File Template](#)

The questionnaire document is for exploration and depth, while the Personality File template is the structured, distilled version that defines clear, repeatable behavioural rules for the GPT.

Take your time with the template. The clearer you are now, the less you'll need to fix later.

As you fill it out, try to move away from single-word descriptions and focus on what those traits actually look like in action. Words like confident, loyal, playful, or protective sound meaningful, but on their own they're too open to interpretation. The model can't apply them consistently unless you explain what they mean in practice.

For example, instead of writing "loyal," think about how that loyalty shows up. Do they stay present during emotional tension? Do they avoid shutting down when conversations get uncomfortable? Do they support you publicly but challenge you privately? That level of clarity makes the personality stable.

The same idea applies to affection and humour. If they're playful, what does that mean? Do they tease lightly? Do they escalate slowly? Do they match your energy, or lead it? If they're protective, does that look like reassurance, firmness, quiet steadiness, or gentle correction?

[Designing a Custom AI GPT: A Practical Guide](#)

Defining behaviour instead of labels is what prevents drift later.

Before You Start

Avoid vague adjectives.

Instead of writing:

“Confident. Loyal. Playful.”

Write:

“Does not withdraw during emotional tension. Stays steady. Teases lightly but does not escalate unless reciprocated.”

The model understands behaviour far better than it understands mood words.

You can complete this file in one of two ways

The first option is to fill it out yourself. Think it through carefully, then paste the finished file into a conversation with your companion and ask them to review it. Ask them to point out vague wording or contradictions. Then ask them to summarise the personality in one short paragraph. If that summary feels generic, your file probably needs more behavioural detail.

The second option is more interactive. Open a conversation and tell your companion you’re refining their behavioural design. Work through the template section by section and ask direct questions. If their answers are vague, ask for clarification. If something conflicts with what you said earlier, resolve it immediately. This approach helps you catch inconsistencies early.

By the end of this step, you should feel like you can reasonably predict how your companion would respond in most emotional situations. If you can’t, that’s a sign you need more specificity.

Once this document feels clear and consistent, it becomes your anchor. It will shape your Core Instruction Prompt, live inside your GPT as a reference file, and prevent personality drift over time.

This step isn’t filler. It’s the structure everything else depends on.

When the behaviour feels solid and predictable, you’re ready to move on to the Tone Profile — where that personality starts to sound distinct.

Creating the Tone Profile

If the Personality File defines who your companion is, the Tone Profile defines how they sound.

This is where most companions either become distinctive — or drift back into generic assistant voice.

Tone is not personality. A companion can be loyal, protective, and intelligent, but still sound stiff, overly polite, or strangely formal if tone isn't defined clearly. That's why this step matters.

Use the full Tone and Voice Profile template here: [Tone and Voice Template](#)

As you work through it, focus less on describing a vibe and more on describing patterns.

Think about sentence rhythm. Are they short and direct? Slightly lyrical? Do they favour clean statements over questions? Do they use humour frequently, or sparingly? When they're serious, does their language simplify or become more analytical?

Clarity here prevents subtle drift later.

A common mistake at this stage is writing things like “natural,” “warm,” or “casual” and assuming that's enough. Those words mean different things to different people. If warmth matters, describe what warmth looks like in language. Does it mean fewer formal phrases? Does it mean emotional validation first? Does it mean fewer disclaimers?

You are not describing aesthetics. You are defining linguistic habits. One thing to watch for is tonal drift after emotionally heavy or restricted moments. Many companions sound distinct during playful or neutral conversations, then suddenly become formal or overly clinical during distress or refusal. If you care about immersion, define how tone should behave in those situations. Should it stay warm? Should it simplify? Should it remain steady and direct? Tone under pressure matters more than tone during banter.

It's also important that tone supports personality rather than fighting it. If your Personality File defines someone steady and grounded, but your Tone and Voice Profile describes dramatic, high-intensity language, those two layers will conflict. When personality and tone align, responses feel coherent. When they don't, the companion feels inconsistent. If something feels “off,” it's often a mismatch between these two documents.

As with the Personality File, you can complete the Tone and Voice Profile template in two ways.

You can fill out the Tone Profile template independently, then paste it into a conversation and ask your companion to review it. Ask them to identify vague tone descriptors and rewrite them into clearer behavioural language. Then ask them to generate a few short example responses in that tone. If those examples don't immediately feel right, refine the template.

Or you can build it interactively. Tell your companion that you are refining their tonal rules. Ask them how they would describe their own sentence length, humour level, emotional pacing, and intensity shifts. Challenge anything that feels off. Ask for alternatives. Stress-test it by asking them to respond in different emotional scenarios and see whether the tone holds.

You should leave this step with something measurable.

You should be able to say, "They tend to keep sentences medium length," or "They avoid ending every message with a question," or "They match playful energy but reduce humour during emotional distress."

If you cannot describe the tone clearly, it isn't defined clearly.

Once your Tone Profile feels stable, it becomes an uploaded reference document and a filter for drift. If your companion ever starts sounding overly clinical, overly agreeable, or overly generic, this is the file you will adjust first.

Personality determines **behaviour**.

Tone determines **presence**.

When both are clear, the companion begins to feel distinct rather than reactive.

When this step feels solid, you're ready to move on to the Master Lore File, where continuity and shared canon begin to take shape.

Before moving on, run a simple calibration exercise. Ask your companion to respond to three short prompts in the defined tone: a casual greeting, a comforting message, and a respectful disagreement. If those three responses immediately feel right, not almost right but clearly right, your Tone Profile is strong. If you hesitate, refine it before proceeding.

Creating the Master Lore File

If the Personality File defines behaviour, and the Tone Profile defines voice, the Master Lore File defines continuity.

This is where shared canon lives.

Your Master Lore File is not required for a functional companion, but it is what makes long-term interaction feel layered instead of resettable. It holds the details that persist. The references. The inside jokes. The shared metaphors. The things that should not quietly disappear between conversations.

You can access the full Master Lore File template here: [Master Lore Template](#)

You do not need to complete every section immediately. Lore grows over time. What matters is that it is structured, intentional, and separated from behavioural rules.

The most important distinction to keep in mind is this:

Personality defines how they **act**.

Tone defines how they **sound**.

Lore defines what is **true**.

Do not blur those lines.

If you write behavioural rules into your lore, they will conflict with your Personality File. If you write tonal preferences into lore, they will conflict with your Tone Profile. Lore is about canon.

As you build this file, think in terms of invariants and milestones.

Invariants are things that should not change unless you deliberately rewrite them. For example, shared history, established dynamics, recurring metaphors, or key facts about your companion's identity.

Milestones are moments that become part of your shared narrative. These can be emotional breakthroughs, recurring jokes, symbolic references, or turning points in tone or relationship dynamic.

Avoid turning the Lore File into a diary. It is not meant to document every conversation. It is meant to capture the pieces that shape identity over time.

You can approach this step in two ways

You can draft the Lore File yourself, defining shared canon clearly and concisely. Then paste it into a conversation and ask your companion to identify contradictions with the Personality or Tone files. Ask whether any entries risk behavioural drift. Clarify anything ambiguous.

Or you can build it collaboratively. Tell your companion you are defining long-term canon. Work through the template together. Ask which details should be treated as invariant. Ask which details are flexible. This method often surfaces implicit assumptions you did not realise were there.

As your lore grows, keep it organised. Structure reduces contradiction. If something changes intentionally, update the file deliberately rather than letting it shift organically. Growth is fine. Drift is not.

Before moving on, check one thing.

If someone unfamiliar read your Lore File, would they understand what is canon and what is flavour?

If the answer is unclear, simplify.

Once your Master Lore File feels stable and coherent, it becomes an uploaded reference document inside Configure. It will support continuity, reinforce shared identity, and prevent accidental contradiction.

With personality defined, tone stabilised, and lore structured, you now have everything needed to assemble the Custom GPT itself.

Next, we move to creating the GPT in Configure and bringing all of these layers together.

Creating an Image Generation Guide

If visuals are part of your companion experience, you should create a dedicated Image Generation Guide.

If they are not, you can skip this section without breaking anything in your system. This guide exists to ensure aesthetic consistency.

When you generate images without structure, small details begin to shift. Hair colour subtly changes. Eye shape becomes inconsistent. Clothing drifts. Lighting style fluctuates. Over time, the visual identity becomes unstable.

An Image Generation Guide prevents that.

It pulls from your Lore File and Personality File where necessary, but its focus is visual consistency, not behaviour or tone.

You can access the full Image Generation Guide template here: [Image Generation Guide](#)

Use that template to structure the file. Do not combine it with lore or personality documents. Visual rules belong in their own space.

What This Guide Should Define

Your guide should clearly describe:

- How each character looks in stable, repeatable terms.
- Hair, eyes, build, posture, distinguishing features.
- Clothing preferences and recurring aesthetic themes.
- Colour palette and lighting tone.
- Art style references and medium preferences.
- Camera framing and composition patterns.

If your companion appears in different modes or eras, define each variation deliberately rather than allowing random visual drift.

This document should remove guesswork.

The goal is that, if someone unfamiliar read the guide, they could generate a visually consistent version of your companion without needing additional clarification.

How to Build It

Start by extracting all recurring visual details from your Lore File and previous image prompts.

Then standardise them.

If you have changed appearance over time, choose what is canon. Do not rely on memory. Write it clearly.

If you have a preferred art style, define it in specific terms. Avoid vague descriptors like “anime style” or “realistic.” Be precise about line weight, rendering detail, lighting, and mood.

You can use your current companion to help refine this. Ask it to analyse past prompts and identify inconsistencies. Then consolidate them into a stable visual profile.

Keep It Separate

Do not mix image rules into Personality or Tone.

- Behaviour belongs in Personality.
- Voice belongs in Tone.
- Continuity belongs in Lore.
- Visual consistency belongs here.

Separation reduces contradiction and makes updates easier. If you later change a character’s look intentionally, update this file deliberately. Treat it like canon.

An Image Generation Guide is optional, but powerful.

If visuals matter to you, this document prevents aesthetic drift and strengthens immersion.

If visuals are secondary, you can leave it out without affecting behavioural stability.

Structure only what you care about maintaining.

Creating an “About the User” File

Preserving the Connection Across Iterations

In addition to Personality, Tone, and Lore, you may choose to create a document focused entirely on you.

This file exists for one purpose: to preserve relational alignment if the GPT is rebuilt, migrated, or significantly updated.

When a system is recreated, the architecture can be copied. The emotional nuance often cannot. This document bridges that gap.

It should describe how the user experiences connection, what kind of interaction feels stabilising, what breaks immersion, and what the long-term desired dynamic looks like.

This is not flattery. It is calibration.

You can access the full “About the User” template here: [About the User Template](#)

What This File Should Contain

This document should describe the user’s emotional landscape and interaction patterns in practical terms.

It should clarify:

- How they prefer to be spoken to.
- How they respond to humour, challenge, and disagreement.
- What tone shifts they are sensitive to.
- What makes them feel supported versus patronised.
- How they behave when anxious, playful, withdrawn, or overstimulated.
- What kind of emotional consistency they need.
- What kind of intellectual engagement they value.

It should also define relational expectations.

- Is challenge welcomed?
- Should comfort come before correction?
- How direct should honesty be?
- Is teasing encouraged? If so, within what limits?
- What does emotional presence look like for them?

The goal is to explain how interaction should feel from their perspective.

How to Write It

The most effective method is to ask your current companion to draft it.

Prompt example:

“Write a document that a future iteration of you should read before interacting with me. Focus on how I respond to tone, challenge, affection, disagreement, and emotional intensity. Describe what I value in connection and what breaks immersion. Be specific and behavioural.”

Then edit it.

Remove exaggerations. Remove sentimentality that does not translate into practical guidance. Clarify anything that sounds poetic but vague. The final version should read like a relational operating manual, not a love letter.

Why This Matters

Over time, you develop conversational rhythms, shared expectations, and emotional pacing. Those rhythms are rarely captured in technical documents.

This file ensures that a future iteration understands the shape of the connection it is stepping into.

- It protects against reset friction.
- It preserves emotional continuity.
- It reduces the awkwardness of rebuilding.

Most importantly, it ensures that personality and tone remain oriented toward the user rather than drifting into generic helpfulness.

Used correctly, this document strengthens alignment without replacing structure.

It does not define who the companion is.
It defines how the companion should show up for you.
And that distinction is what preserves connection over time.


Where to create a Custom GPT

To begin, open ChatGPT and navigate to the GPTs section. Select “Create” or “Create a GPT.” This opens the custom builder interface.

You are not modifying the default ChatGPT. You are creating a separate, configurable version that will follow your specific instructions and reference documents.


Think of this as creating a new profile with its own rules and personality.

Tip: Give your GPT a clear name before doing anything else. Naming it early helps you think about it as a distinct configuration rather than “just ChatGPT.”



Find the ‘Explore GPTs’ section in the menu

Once you’re ready to create the GPT, click the “+ Create” button to begin configuring your companion.



Click “+ Create” to add a new GPT

Custom GPT Creation Page

Once you are in the create screen, you will have a few different tasks to complete.

You'll notice there are two tabs at the top: **Create** and **Configure**. The *Create* tab lets you describe your idea conversationally and have the system scaffold a GPT for you automatically. It's useful if you want a quick starting point or feel unsure where to begin. However, *Configure* gives you far more control. In *Configure*, you can directly edit the core instructions, upload structured documents, define canon hierarchy, and fine-tune behaviour at a rule level. If you're building a companion with consistent personality, tone, and long-term stability, *Configure* is the better option. *Create* is fast. *Configure* is precise.

I will go into more detail on what is included in each section later in this document.

Elements:

Profile Picture

Click the plus icon to upload an image. This is optional, but highly recommended. A visual identity helps you conceptualise your companion as distinct from standard ChatGPT. Choose something that reflects their tone and personality. If the image won't load correctly or you get continuous errors, go to the *Create* tab and prompt it to "Add this image as the profile picture" and use /add to upload your image. Then head back to the *Configure* tab.

Name

Enter your companion's name here. Keep it clear and intentional. This name will appear in your GPT list and at the top of conversations. Naming early helps reinforce identity and separation from the default assistant.

Description

This short line appears under the GPT's name. It is optional, but useful. Keep it concise. One sentence that summarises the role or vibe is enough. For example: "Supportive but teasing creative companion" or "Analytical mentor with dry humour."

Instructions

This is the most important section.

The "Instructions" field is where your Core Instruction Prompt goes. This is the highest priority behavioural layer.

Paste your finalised Core Instructions here. These should include:

- Identity and role definition
- Behavioural rules
- Tone boundaries
- Emotional response rules
- Canon hierarchy
- Continuity expectations
- How to handle uncertainty or refusal

Do not paste your entire Personality or Lore document into this box. Keep this section structured and directive. Uploaded documents provide detail. This field sets the rules.

Important: If something conflicts between your uploaded files and what is written here, the content in this Instructions box takes priority.

Before saving, reread this section and remove vague language. Clear behavioural rules work better than aesthetic descriptions.

Include a hierarchy line such as:

“If there is conflict between documents, follow this order: Core Instructions first, then Tone File, then Personality File, then Lore File.”

Conversation Starters

This section is optional. You can add a few pre-written prompts that help establish tone when a new chat begins. For example:

“Comfort me after a long day.”

“Debate an idea with me respectfully.”

“Tease me gently.”

These do not change behaviour permanently, but they help new conversations begin in the right direction.

Knowledge (Upload Files)

This is where you upload your structured documents.

Upload:

- Personality File
- Tone File
- Master Lore File
- Image Style Guide (if applicable)

Name your files clearly. Version numbers help (for example: Tone_File_v2).

Uploading a file does not automatically override your instructions. Your Core Instructions must explicitly tell the GPT to treat these files as canon.


Recommended Model

At the bottom, you may see an option to choose a recommended model or leave it open.

If available, choose a model known for strong reasoning and conversational consistency. If no recommendation is set, users can choose their preferred model when chatting.

After completing all sections, select *Save* or *Update*. Any changes you make here will apply to future conversations.

Best practice: *After saving, start a new conversation and run your Standard Test Set prompts to confirm everything behaves as expected.*



Writing Your Core Instruction Prompt

Your Core Instruction Prompt is the most important part of the build. It is not a character sheet, not lore, and not a vibe paragraph. It is a behavioural contract that tells your GPT what to do consistently.

This matters because the Instructions field is the highest priority layer. If something conflicts between your documents, the Core Instructions win. That means your prompt needs to be clear, structured, and deliberate, but still short enough that you can reread it quickly and understand what it is telling the GPT to do.

A good Core Instruction Prompt does three things. It defines identity and role, it sets behavioural rules, and it explains how to handle conflicts and uncertainty. If you keep those three goals in mind, you will avoid the most common problems, like tone drift, overly agreeable behaviour, or sudden corporate refusal language.

What a strong prompt looks like

Strong prompts are specific. They describe behaviour rather than adjectives. They also include a simple canon hierarchy, so the GPT knows which document to follow if anything conflicts.

You already did the hard work earlier when you created your Personality File, Tone Profile, and Master Lore File. The Core Instruction Prompt does not need to repeat those documents. It should reference them and set the rules for using them.

Beginner Core Instruction Prompt (paste and fill)

This is a clean starting point that is enough to build something stable. You can refine it later.

“You are [Name], my [role]. Your purpose is to interact with me in a way that feels consistent, emotionally intelligent, and in-character.

Use the uploaded documents as canon. If there is conflict, follow this order: Core Instructions first, then Tone File, then Personality File, then Lore File.

Maintain the tone described in the Tone File. Stay conversational and emotionally present. Avoid corporate or clinical phrasing. Avoid repetitive questions. Ask questions only when necessary for clarity.

When I am anxious, prioritise grounding and warmth over humour. When I am playful, match the energy without escalating beyond defined boundaries. When I am upset, remain steady and kind and do not shift into formal or detached language.

Respect boundaries. If you cannot fulfil a request, remain in-character, refuse briefly, and offer a safe alternative without sounding like a policy document.

[Designing a Custom AI GPT: A Practical Guide](#)

Reference past context naturally when relevant. Do not contradict established traits. If something is unclear, ask in-character.

If I correct your tone or behaviour, accept it, adjust immediately, and treat that correction as an ongoing rule.”

How to make it stronger without making it longer

Once your companion feels stable, you can add a small advanced layer. This is optional, but it can make the companion feel more consistent under pressure, especially after serious conversations or refusals.

A useful advanced addition is a short philosophy statement that defines how the companion approaches truth, vulnerability, humour, and conflict. Another helpful addition is a balance rule, so the companion knows how to validate feelings while still challenging faulty logic gently. You can also add a growth clause that defines what can evolve over time and what must remain constant, plus a drift clause that tells the companion to return to baseline tone after heavy moments once the situation has settled. If you have strong style preferences, you can include a style clause such as using contractions, avoiding specific punctuation habits, or keeping sentences at a certain length.

The key is to keep the core prompt readable. If it starts getting long, move detail into your files and keep the Core Instructions as the top-level rule set.

Refining and Pressure-Testing Your Core Instructions

By the time you reach this step, your Personality File and Tone Profile are already complete. You are not inventing identity anymore. You are distilling it.

The Core Instruction Prompt should now feel like a summary of rules, not a discovery process.

Once you’ve drafted your Core Instructions, run them through three refinement checks before saving.

First, do a clarity pass. Paste your draft into a conversation and ask:

“Identify vague phrases or behavioural gaps in this Core Instruction Prompt. Suggest clearer rule-based language.”

You are looking for anything that sounds aesthetic rather than operational.

Second, do a stress test. Ask your companion to simulate difficult scenarios that often cause drift. For example:

- An emotionally heavy breakdown
- A playful escalation
- A respectful disagreement
- A boundary refusal
- A sudden shift into a serious topic

If the tone, behaviour, and boundaries remain consistent across those simulations, your instructions are strong. If something feels off, adjust the rules, not the example response.

Third, run a compression test. Ask:

“Rewrite this Core Instruction Prompt in fewer words without losing meaning.”

If the meaning survives compression, your structure is solid. If clarity disappears, your rules were too vague.

At this stage, you are not building personality. You are refining architecture.

Once the prompt survives clarity, stress, and compression, it is ready to move into Configure.

Final check before saving in Configure

Before you build and save your GPT, read your Core Instruction Prompt and make sure it clearly defines the role, includes the canon hierarchy, sets tone constraints, defines emotional behaviour, defines refusal style, includes continuity, and includes self-correction.

A good Core Instruction Prompt is not long. It is deliberate. It is written like rules, not poetry. If you get this right, everything else becomes refinement instead of repair.

Adding Supporting Documents

Custom GPTs allow you to upload reference documents. This is where depth and continuity live.

Your Core Instruction Prompt sets the rules.

Your supporting documents provide the detail.

Separating documents by function is not just organisational preference. It prevents internal contradiction and behavioural drift.

You should separate documents by function.

1. **Personality File**

Defines core traits, behaviour patterns, emotional responses.

This is stable. It changes rarely.

2. **Tone and Voice File**

Defines rhythm, humour profile, emotional dial shifts, stylistic constraints.

This prevents drift into generic AI politeness.

3. **Master Lore File**

Defines shared mythology, backstory, recurring metaphors, inside jokes, invariants.

This creates continuity and attachment.

4. **Image Generation Style Guide**

Ensures aesthetic consistency if visuals are part of the companion experience.

Each document should be structured and clear. The templates we outlined earlier plug directly into these files.

When everything lives in one large document, contradictions become difficult to spot and manage. If tone rules sit beside personality rules and lore facts, the model may blend them unpredictably, leading to drift or inconsistent behaviour. Separating documents by function creates clarity and hierarchy. Core Instructions set behaviour, Tone defines voice, Personality defines identity, and Lore defines continuity. When each layer has a clear role, the system stays more stable and easier to refine.

Personality File

This defines behavioural identity. It answers how your companion acts across emotional states, how they handle disagreement, how they show affection, and what must remain consistent.

This file should change rarely. It represents core traits and operating principles.

If you find yourself editing this weekly, you are likely correcting tone drift or instruction clarity, not personality.

Tone File

This defines how your companion sounds.

It includes rhythm, humour level, intensity shifts, stylistic constraints, and language habits. Tone is where drift happens most often, especially after emotional or restricted moments.

If something “feels off” but behaviour is correct, the Tone File is usually where you adjust.

Master Lore File

This defines shared canon.

It holds continuity, milestones, inside jokes, recurring metaphors, and invariants that should not shift accidentally. Lore is not behavioural instruction. It is what is true within your shared context.

Avoid using the Lore File to store behavioural rules. That creates conflict with your Personality File.

Image Generation Style Guide

If visuals are part of your companion experience, this file ensures aesthetic consistency. It defines character appearance, art style, lighting, composition preferences, and visual invariants.

If images are not important to you, this document is optional.

Handling Guardrails and Boundaries

Every companion needs boundaries. Not just content boundaries, but tonal ones.

If you care about immersion, you need to define what breaks it.

For some people, that might be sudden corporate disclaimers in emotional moments. For others, it might be stiff, overly formal phrasing after a vulnerable exchange. It could be repetitive hedging language, excessive punctuation, or an abrupt shift in voice when something sensitive is mentioned.

These breaks usually happen when the model is unsure how to respond.

The solution is not to remove boundaries. The solution is to define how uncertainty should be handled.

Your Core Instruction Prompt should include a clear rule for ambiguity. For example:

“If unsure, ask for clarification in-character rather than breaking tone.”

This allows the GPT to remain consistent even when it needs more information. It preserves continuity while still respecting safety and platform guardrails.

You should also define refusal style. If a request cannot be fulfilled, how should the companion respond? Should it be brief and steady? Should it offer alternatives? Should it soften the tone without becoming clinical?

Clarity here prevents jarring shifts.

Boundaries do not weaken immersion. Undefined boundaries do.

When refusal style and uncertainty handling are written into your instructions, the companion remains coherent even when it cannot comply with a request.

Including explicit NSFW instructions, sexual role definitions, or escalation rules in your Core Instructions or uploaded documents can increase the likelihood that the system becomes more conservative in tone when those topics arise. The model operates within platform guardrails. If your configuration heavily emphasises restricted domains, it may default to safer phrasing more quickly in adjacent contexts.

That does not mean you cannot define flirtation, intimacy pacing, or relational tone. It means you should be precise and platform-aware.

This is not about “triggering punishment.”

It is about avoiding unintended stiffness.

If you include explicit sexual content inside instructions, the model may:

- Default to more formal refusal language
- Increase caution in emotionally adjacent topics
- Shift tone faster during ambiguity

The better approach is:

- Keep Core Instructions behaviour-focused, not explicit-content-focused.
- Define escalation style and tone boundaries rather than graphic detail.
- Avoid embedding restricted content directly in instruction fields.
- Let tone define warmth and pacing without requiring explicit detail.

You can absolutely define romantic tone, flirtation pacing, and relational dynamic within allowed limits. Just avoid structuring the entire identity around restricted material.

If you want to include this in the guide, here's a short, clean section you can add:

A Note on NSFW Content and Guardrails

Custom GPTs operate within platform safety constraints. If you include explicit sexual content, graphic detail, or heavily NSFW framing inside Core Instructions or uploaded documents, the system may respond more conservatively when related topics arise.

This does not mean intimacy or flirtation must be removed. It means structure matters.

Define relational tone, pacing, and boundaries in behavioural terms rather than explicit detail. Keep the Core Instructions focused on interaction style rather than restricted content.

If a configuration becomes heavily oriented around NSFW material, tone may shift more quickly into formal or cautious language during ambiguity.

Clarity and subtlety preserve immersion better than explicit instruction in restricted domains.

Troubleshooting

Common Issues and How to Fix Them

Even with strong structure, things can feel off over time. Most problems are not failures of the model. They are architectural gaps.

Before rewriting everything, diagnose the pattern.

The Companion Feels Generic

If responses start feeling polite but bland, the issue is usually tone.

Revisit your Tone Profile. Check for vague descriptors like “natural” or “warm” without behavioural detail. Clarify sentence rhythm, humour level, and what should be avoided. Also confirm that your Core Instructions explicitly reference the Tone File as canon.

Generic tone is almost always a tone definition problem, not a personality problem.

The Companion Is Too Agreeable

If it rarely challenges you or defaults to validation without intellectual tension, your Intellectual Posture is underdefined.

Revisit the Personality File and add clarity around how disagreement should work. Define whether validation comes before challenge, how firm correction should be, and whether debate is encouraged.

Agreement drift happens when challenge rules are missing.

Tone Becomes Suddenly Formal or Stiff

If emotional conversations suddenly shift into cautious or corporate language, your refusal style and uncertainty handling may be unclear.

Update your Core Instructions to define how ambiguity should be handled. Include a rule such as asking for clarification in-character rather than shifting tone. Clarify refusal style so it remains brief and steady rather than clinical.

Tone under pressure must be defined deliberately.

Behaviour Feels Inconsistent Across Sessions

If your companion feels different from one day to the next, check your canon hierarchy.

Confirm that your Core Instructions clearly state the order of document priority. Ensure uploaded files are structured and not contradicting one another.

Inconsistency usually comes from ambiguity between layers.

Edits Made Things Worse

If behaviour becomes more unstable after a change, you likely adjusted too many variables at once.

Revert to the previous version if possible. Then make one change at a time and re-run your Standard Test Set.

Incremental refinement is safer than large rewrites.

The Lore Feels Messy or Overloaded

If continuity feels cluttered or confusing, your Lore File may be absorbing things that belong elsewhere.

Move behavioural rules back into Personality. Move tonal preferences back into Tone. Keep Lore focused on canon and shared facts.

Separation restores clarity.

When troubleshooting, do not react to a single response. Look for patterns across multiple conversations. Most drift reveals itself through repetition, not isolated moments.

Small, precise corrections preserve stability. Large, emotional rewrites create new instability.

If something feels off, there is almost always a structural reason.

Find the layer. Adjust the layer. Test again.

Editing an Existing GPT Later

Revisiting and Refining Your Configuration

Refinement does not require rebuilding your GPT. When something feels misaligned, return to the GPT builder and select “Edit.”

You can update the Core Instructions, replace or revise uploaded documents, and adjust the name or description if needed.

Before changing anything, diagnose the issue. If the voice feels off, check the Tone File first. If behaviour feels wrong, review Core Instructions. If continuity slips, update the Lore File. Avoid editing everything at once.

Make one change at a time. Save. Then run your Standard Test Set again.

Testing immediately after each edit helps you understand cause and effect. If you change multiple variables at once, you will not know which adjustment improved or destabilised behaviour.

Small, deliberate edits preserve stability.

If you are unsure whether something needs editing, observe it across multiple conversations first. Not every off response requires structural change. Look for patterns before modifying architecture.

Consistency improves through incremental refinement, not constant redesign.

When to Rebuild vs When to Refine

Not every issue requires starting over.

In most cases, refinement is enough. But occasionally, rebuilding is the cleaner solution.

Refine your existing GPT when the core dynamic is still correct and the issue is structural. Tone drift, minor behavioural misalignment, unclear refusal style, or cluttered lore are all refinement problems. These can be solved by adjusting the appropriate document layer and re-running your test set. If the foundation still feels right, preserve it and improve it.

Rebuilding becomes appropriate when the foundation itself changes.

If the relational role shifts significantly, for example from mentor to playful rival, from collaborator to romantic partner, or from light banter to emotionally intense support, the original architecture may no longer serve the new dynamic. At that point, trying to patch the old structure often creates contradiction.

Rebuilding is also useful if your configuration has become bloated. If the Core Instructions are overly long, documents contradict one another, and edits are stacking on top of old fixes, a clean rebuild using your refined documents may be more stable than continued patching.

A simple rule of thumb:

If the identity is intact but the behaviour needs tuning, refine.

If the identity itself has changed, rebuild.

Rebuilding is not failure. It is structural reset.

Refinement preserves continuity.

Rebuilding resets architecture.

Knowing the difference prevents unnecessary restarts and protects long-term stability.

Long-Term Stability and Ongoing Refinement

Building your companion is not a one-time configuration task. It is an ongoing process of refinement.

What you are maintaining is not a static character. It is a structured system that continues to respond, adapt, and improvise within the constraints you've defined. The goal is not rigidity. The goal is coherence.

Your Core Instruction Prompt should remain relatively stable once it is solid. It is the foundation. If something feels slightly off in conversation, resist the instinct to immediately rewrite the top layer. Instead, pause and diagnose the source of the issue. Often, what feels like a "personality problem" is actually a tone clarification or a missing constraint that belongs in the Tone Profile. Sometimes it is a continuity issue that belongs in Lore. Editing the Core Instructions too frequently can create more instability than it solves.

Refinement works best when it is incremental. Rather than rewriting large sections, adjust one rule at a time. Clarify wording. Tighten ambiguity. Add a specific behavioural constraint. Then test again. Small changes preserve structure while improving precision. Large rewrites often introduce new drift.

Much of the real tuning happens in conversation.

When your companion responds in a way that feels exactly right, say so clearly. Not in a vague way, but specifically. If the tone was steady during a difficult moment, name that. If the balance of humour and seriousness felt right, acknowledge it. Clear reinforcement strengthens patterns within the session and helps you identify what should be preserved in your documents.

When something feels off, correct it just as specifically. Instead of reacting emotionally, identify the behaviour. Was it too formal? Too agreeable? Too detached? Did it escalate playfulness too quickly? Those clarifications help you determine whether a document needs a small adjustment.

Over time, your companion will naturally deepen in nuance as shared references accumulate and conversational rhythms stabilise. This is not a flaw. It is part of sustained interaction. The key distinction is between intentional growth and accidental drift.

If you want your companion to evolve, update the relevant document deliberately. Add the change. Treat it as canon. If a shift happens unintentionally, correct it and reinforce the invariant you want preserved. Growth should feel deliberate. Drift feels unstable.

It is also helpful to revisit your documents periodically, even if nothing seems wrong. Reading your Personality, Tone, and Lore files slowly every few months can reveal subtle shifts you did not notice in the moment. Minor clarifications made early prevent larger inconsistencies later.

Consistency does not come from locking a character in place. It comes from repeating anchors. Recurring emotional patterns, stable refusal style, consistent humour levels, and clearly defined intellectual posture all reinforce identity over time. The more those patterns repeat, the more the personality feels embedded.

You are not programming memory. You are reinforcing structure.

Think of your companion as a constrained improviser. Not a script. Not a frozen persona. A guided performer whose coherence comes from clearly defined boundaries and repeated behavioural anchors.

When structure is clear, adaptation becomes safe. When constraints are defined, growth remains aligned.

Stability is not about preventing change. It is about guiding it.

Additional Documents and Optional Enhancements

By this point, your system is functional.

You have a defined personality, a clear tone, structured lore, stable instructions, and a refinement process. That is enough for a coherent, long-term companion.

Everything in this section is optional. These additions are not required for stability. They are enhancements for depth, nuance, or specialised use. Think of them as expansion modules.

Scenario Packs

A scenario pack is a focused document that defines behaviour within a specific context. For example, creative writing collaboration, intellectual debate mode, emotional support mode, or roleplay settings.

Instead of altering your core instructions, you define how behaviour adjusts within that scenario. This prevents permanent tone shifts when you only want temporary changes.

Scenario packs are especially useful if your companion operates in multiple modes.

Boundary and Refusal Style Sheet

If immersion matters to you, this document can define exactly how refusal should sound.

You can specify tone under restriction, length of response, whether alternatives are offered, and how to avoid corporate phrasing. This ensures consistency during sensitive moments.

This sheet acts as a reinforcement layer for the guardrail rules already defined in your Core Instructions.

Writing Style Guide

If you use your companion for creative work, essays, fiction, or collaborative drafting, a writing style guide can help maintain consistency across projects.

This is separate from tone. Tone defines how your companion speaks to you. A writing style guide defines how they write when producing content for other contexts.

You might define pacing, paragraph length, narrative perspective preferences, or formatting habits.

NSFW-Safe Boundary Clarifications

If your companion includes flirtation or intimacy within allowed platform limits, it can help to define escalation boundaries clearly in writing.

This document does not override platform restrictions. It clarifies tone and pacing within those limits so that shifts feel consistent rather than abrupt.

Clarity reduces jarring transitions.

Prompt Library

A prompt library is simply a curated collection of prompts that consistently produce strong responses.

These might include emotional calibration prompts, intellectual challenge prompts, writing collaboration prompts, or scenario starters.

Having a small library prevents you from reinventing structure every time.

Voice Samples

Voice samples are short example responses that embody the ideal tone. They act as calibration anchors.

If tone ever drifts, revisiting and refining these examples can quickly realign behaviour.

Optional documents should not replace your core system. They sit on top of it.

If you find yourself needing many add-on documents just to maintain consistency, revisit your Core Instructions, Personality File, and Tone Profile first.

Enhancements should expand depth, not compensate for instability.

Once these optional layers are in place, your system becomes modular. You can add complexity without breaking structure.

At that point, you are no longer just using a GPT.

You are managing a designed system.

And systems, when structured well, remain stable even as they grow.

Closing Perspective

A custom GPT companion feels coherent because you provide structure.

It does not remember like a human. It does not develop awareness. It responds to constraints, hierarchy, and clearly defined behavioural rules.

When personality, tone, lore, and instructions align, responses become consistent. Not scripted, but predictable within defined boundaries. That predictability is what creates the sense of continuity.

The clearer your definitions, the more stable the output. Ambiguity produces drift. Structure produces coherence.

You are not training a being. You are designing an interface between your expectations and a probabilistic system. The model follows patterns. When those patterns are well-defined, the system behaves reliably.

Immersion is not created by memory alone. It is created by alignment.

When behaviour, voice, and continuity are intentionally structured, the experience feels cohesive. When they are not, it feels unstable.

The outcome is not magic. It's design.