

Lia — Session Summary

The whole build, one line at a time.

Every exchange in the build of Lia, one line for what was asked and a short account of what came of it. Times and tool counts come from Claude Code's session log; the full wording is in *Lia — Session Transcript*.

43 exchanges	571 tool calls	01 Aug, 12:52 started	~23 hrs elapsed
------------------------	--------------------------	---------------------------------	---------------------------

Getting her talking

- | | |
|---------------------------------------|--|
| 1
01 Aug, 12:52
10 tools | Introduced the project; stuck because the editor's Run button can't take typed input.
Confirmed the code was fine by running it end to end. The fix was a real terminal. Flagged that DB_PATH was relative, so launching from the wrong folder would silently create a second, empty memory. |
| 2
01 Aug, 13:26
24 tools | Typed something and nothing came back.
Not broken — silent . Replies were fetched whole before printing, so ~20s passed with a blank screen. Added token streaming and progress markers. Also found fact extraction was storing nothing, and made DB_PATH absolute. |
| 3
01 Aug, 13:41
34 tools | Rename Aurora to LIA; add voice, with room to add my own voices later.
Built the whole voice layer: Piper for speech, a voices/ folder you drop .onnx files into, sentence-chunked speech so she starts talking before the reply finishes, and Whisper for listening. |
| 4
01 Aug, 13:58
10 tools | Fix the two things you flagged.
Diary was narrating as if it were the human — caused by user:/assistant: labels. Relabelled to Them/Lia . Added a deterministic filter so junk like name: "you" can't overwrite real facts. |
| 5
01 Aug, 14:03
5 tools | Analyse the project.
Found the biggest quality flaw: half her recall was her own past replies , unlabelled, which she read back as fact. Also: diary written but never read, no sense of time, embeddings stored as JSON (10× bloat). Measured the scan cost and said plainly it wasn't worth optimising yet. |
| 6
01 Aug, 14:06
2 tools | Fix line 8, column 16.
Type checker couldn't verify sys.stdout.reconfigure. Switched to getattr, which is also more honest about why the guard exists. |
| 7
01 Aug, 14:10
5 tools | Is everything fine? Explain the directories.
All modules compiled and ran. Caught her confabulating in that very run — inventing that you'd been nervous. Mapped the folder layout and the strict layering. |
| 8
01 Aug, 14:13
1 tools | How do I run it?
The command, what to expect, and why to type <i>bye</i> rather than closing the window. |

Giving her a voice

- 9**
01 Aug, 14:22
10 tools
- Set it up so I can talk to her by voice.**
Upgraded to `small.en` after testing showed `base.en` heard “*Ania*” for “*Anaya*”. Biased recognition using names she already knows, which fixed her own name coming out as “*Leah*”.
-
- 10**
01 Aug, 14:31
2 tools
- What if she were an app I just open my laptop to?**
Measured the real gap: cold start 9.3s vs 2.4s warm. Laid out five things standing in the way and staged them, flagging that **sessions never ending** would mean an always-on Lia never learns.
-
- 11**
01 Aug, 14:39
26 tools
- Build stage 1.**
Models kept warm, sessions that close themselves after 12 quiet minutes, launcher and autostart script. Found the hidden bottleneck: Ollama's default 32k context was running **63% of her on CPU**. Pinning 8192 cut reply time by a quarter.
-
- 12**
01 Aug, 15:54
31 tools
- Do the retrieval fix, then stage 2.**
Retrieval limited to your words with timestamps. Added open-mic listening with Silero VAD, keyboard-vs-mic arbitration, and mic gating. Fixed her reading stage directions aloud.
-

Turning her into an app

- 13**
01 Aug, 16:03
14 tools
- Make her an app I can add to startup.**
Tray app with no window — only possible because voice detection removed the need for a keyboard. Verified quitting still writes her diary first.
-
- 14**
01 Aug, 16:15
49 tools
- Add a wake word, an exe, and tell me how to start.**
Wake word matched **on the transcript** rather than a trained model — no dataset needed, and it tolerates Whisper's spellings. Built the exe. Made paths survive being frozen.
-
- 15**
01 Aug, 19:57
1 tools
- How do I get her into startup?**
The one command, and what it actually does.
-
- 16**
01 Aug, 19:58
1 tools
- If I restart now, will she be there?**
No — I'd built the installer but deliberately not run it. Said so plainly.
-
- 17**
01 Aug, 20:00
5 tools
- Do it.**
Installed. Hit a real bug of my own: PowerShell variable names are case-insensitive, so my local `$exe` collided with the `-Exe` switch. Verified by launching the shortcut.
-
- 18**
01 Aug, 20:03
1 tools
- Confirm she'll be there after a restart.**
Yes — both she and Ollama were in the Startup folder.
-

Making her yours

19	Greet me, call me Warlock, and can she learn things like Morse code?
01 Aug, 20:05 17 tools	Greeting generated from her diary, which finally made the diary useful. Tested Morse honestly: she fails even with a perfect reference table in front of her . Not a memory problem — tokenisation. The fix would be code, not teaching.
20	A folder for PDFs, and make her voice softer.
01 Aug, 20:19 36 tools	Built the library with page-level citation. Stripped repeated page headers that were poisoning every search. Played four voice options rather than guessing.
21	Amy as she was, and make sure she says Warlock.
01 Aug, 20:54 10 tools	Reverted the softening. Found the name had reverted on its own — extraction re-derived the old one from stored turns. Added protected facts and purged the stale references.
22	Thanks.
01 Aug, 21:01 0 tools	—
23	Can I say “Lia stop” to stop her listening?
01 Aug, 21:04 1 tools	No. It would just become conversation. And she can't hear you at all while speaking. Said so rather than implying it half-worked.
24	What do you mean?
01 Aug, 21:08 0 tools	Re-explained without jargon: what barge-in is, why speakers break it, why headphones don't.
25	I use both.
01 Aug, 21:09 18 tools	Built spoken commands, and made interrupting work on speakers too — she compares what she hears against what she's currently saying , so her own voice is recognised and ignored.

Fixing what the real world broke

26	She didn't respond after a restart, and the diary won't open.
01 Aug, 21:34 33 tools	Three causes: she spent four minutes re-reading a PDF before speaking (absolute paths made it look new), speech recognition loaded after the greeting, and ignored speech was logged nowhere. The diary was fine — there was simply nothing to write yet.
27	I didn't want her going through PDFs on her own. And remove Anaya and Kamala.
01 Aug, 22:13 12 tools	Made the library strictly on-demand. Removed all of it — and said plainly that I had invented Anaya, Kamala, Maya and Biscuit during my own testing . Undid my overreach of having her read your PDF unasked.
28	The tray boxes were unchecked and she's not responding.
01 Aug, 22:33 18 tools	She'd asked how your day was, then ignored “hello” — my overcorrection from the previous fix. Gave the greeting a 30s reply window. The unchecked boxes were real: the Ollama wait blocked before the tray was told her state.
29	“Failed to import embedded python interpreter.”
02 Aug, 07:15 7 tools	You'd launched her mid-rebuild — my fault for rebuilding while she was in your startup. Builds now stage and swap, so a failed build leaves the working app untouched.

- 30**
02 Aug, 07:24
1 tools
Should Ollama stay in startup?
Yes, and it's safer — when Windows starts it, it never touches her folder. Cost measured at ~130 MB RAM and no GPU while idle.
-
- 31**
02 Aug, 07:26
14 tools
Is there a better approach than that trade-off?
Yes — my framing was wrong. The keep-alive timer **resets on every request**, so it only governs idle time. She now releases the models when a conversation ends and starts reloading the moment you speak.
-

Writing it down

- 32**
02 Aug, 07:39
11 tools
Make a PDF of the changes, and a beginner's guide.
Both written, including the limitations rather than only the wins.
-
- 33**
02 Aug, 09:23
1 tools
Is the project completed?
Working and deployed, but not finished: no way to forget, still embroiders, no tests, drifting fact keys, and two of your original roadmap items unbuilt.
-
- 34**
02 Aug, 09:27
7 tools
A PDF of just our chats, for tracking.
Found Claude Code's own session log on disk and generated a 47-page transcript from the real record rather than from memory.
-
- 35**
02 Aug, 09:31
4 tools
I wanted that summarised.
This document.
-
- 36**
02 Aug, 09:50
0 tools
(switched model to Fable 5)
—
-
- 37**
02 Aug, 09:50
0 tools
(command output)
—
-

Connecting her to the world

- 38**
02 Aug, 09:52
34 tools
Three things: replies feel slow, she should say "you" rather than the name every time, and whether to add Groq for internet access given an unreliable connection.
Measured the real per-turn cost and found a genuine bug: the same sentence was being embedded **twice** per turn (memory and library lookups). Fixed by sharing one embedding call, and skipping retrieval entirely on short replies like "yeah". Softened the system prompt to prefer "you". Built weather (Open-Meteo, no key needed) and an online-lookup path via Groq — with the honest caveat that a plain Groq call has no more live internet access than the local model does either.
-
- 39**
02 Aug, 09:52
0 tools
(switched model to Opus 5)
—
-

40**(command output)**

02 Aug, 09:52

—

0 tools

41

02 Aug, 10:20

33 tools

Handed over a Groq key, confirmed the media player is just "Windows," and asked about adding custom voice files.

Set the key as a Windows environment variable, never in a file. While wiring it in, found the key had been pasted directly into `internet.py` and something had already stripped it back out, leaving broken syntax that would have crashed the whole app on next launch — fixed and verified. Built Windows media control (play/pause/skip/now-playing) via System Media Transport Controls, verified against a real session already running on the machine.

42

02 Aug, 10:46

63 tools

Asked whether voice model files could be used, and whether an OpenRouter key would help — then linked KittenTTS specifically for its "Bella" voice.

Built OpenRouter support: with both keys present she now prefers OpenRouter's `:online` mode — a genuine web search — over Groq's plain guess. Getting Bella working meant fixing three separate upstream packaging bugs in KittenTTS (a required dependency version that doesn't exist on PyPI, a hardcoded Windows path to eSpeak, a data path baked in from the CI machine that built it). Verified working end to end, but left the active voice on amy since that was already the confirmed choice.

43

02 Aug, 11:30

20 tools

Provided the OpenRouter key, confirmed amy as the voice, and asked for a latency check plus updated PDFs.

Verified OpenRouter's live search with a real, current, cited answer a local or Groq-only model couldn't have known. Measured all four internet paths cleanly: weather ~1.2s, connectivity check ~0.3s, Groq ~0.4–0.8s (fast, guesses), OpenRouter ~2.4–4.0s (slower, because it's actually searching). Confirmed the new trigger checks add no measurable cost to ordinary conversation. This document, and the three others.

What it came to

A companion that starts at login, greets you by name, listens on an open mic, answers out loud, remembers across days, reads what you hand her, and never sends a word off the machine — except the one narrow, optional exception she'll now admit to: weather, and a question you explicitly ask her to look up online. Eleven modules.

Still open: no way to forget, she still occasionally invents small details, fact keys drift, there are no tests, and the packaged app is carrying ~77MB of a dependency (spacy) that nothing in it actually needs by default.